

TRPO & PPO

Samsung AI Expert

Jul 10, 2025

Giho Kim

TRPO - Review

TRPO-Review : 1) Surrogate Objective

$$\max_{\pi} J(\pi) - J(\pi_{old}) = \sum_t \mathbb{E}_{s_t \sim p^{\pi}} [\mathbb{E}_{a_t \sim \pi(\cdot|s_t)} [\gamma^t A^{\pi_{old}}(s_t, a_t)]]$$

What is $A^{\pi}(s_t, a_t)$?
 $Q^{\pi}(s_t, \mathbf{a}_t) - V^{\pi}(s_t)$

TRPO-Review : 1) Surrogate Objective

$$\begin{aligned}\max_{\pi} J(\pi) - J(\pi_{old}) &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi(\cdot|s_t)} \left[\gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \\ &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi_{old}(\cdot|s_t)} \left[\frac{\pi(a_t|s_t)}{\pi_{old}(a_t|s_t)} \gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \quad (\text{importance sampling})\end{aligned}$$

What is importance sampling?

$$\mathbb{E}_{p(x)}[f(x)] = \int_x p(x) f(x) dx = \int_x q(x) \frac{p(x)}{q(x)} f(x) dx = \mathbb{E}_{\mathbf{q}(\mathbf{x})} \left[\frac{p(x)}{q(x)} f(x) \right]$$

TRPO-Review : 1) Surrogate Objective

$$\begin{aligned}\max_{\pi} J(\pi) - J(\pi_{old}) &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi(\cdot|s_t)} \left[\gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \\ &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi_{old}(\cdot|s_t)} \left[\frac{\pi(a_t|s_t)}{\pi_{old}(a_t|s_t)} \gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \quad (\text{importance sampling}) \\ &\geq L_{\pi_{old}}(\pi) - C \max_s D_{KL}(\pi_{old}(\cdot|s) \parallel \pi(\cdot|s)) \quad (\text{surrogate objective})\end{aligned}$$


$$L_{\pi_{old}}(\pi) := \sum_t \mathbb{E}_{s_t \sim p^{\pi_{old}}} \left[\mathbb{E}_{a_t \sim \pi_{old}(\cdot|s_t)} \left[\frac{\pi(a_t|s_t)}{\pi_{old}(a_t|s_t)} \gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right]$$

TRPO-Review : 1) Surrogate Objective

$$\begin{aligned}
 \max_{\pi} J(\pi) - J(\pi_{old}) &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi(\cdot|s_t)} \left[\gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \\
 &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi_{old}(\cdot|s_t)} \left[\frac{\pi(a_t|s_t)}{\pi_{old}(a_t|s_t)} \gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \quad (\text{importance sampling}) \\
 &\geq L_{\pi_{old}}(\pi) - C \max_s D_{KL}(\pi_{old}(\cdot|s) \parallel \pi(\cdot|s)) \quad (\text{surrogate objective})
 \end{aligned}$$

Requirements?

- $M(\theta|\theta_i) \leq f(\theta) \quad \forall \theta$
- $M(\theta_i) = f(\theta_i)$



TRPO-Review : 1) Surrogate Objective

$$\begin{aligned}\max_{\pi} J(\pi) - J(\pi_{old}) &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi(\cdot|s_t)} \left[\gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \\ &= \sum_t \mathbb{E}_{s_t \sim p^{\pi}} \left[\mathbb{E}_{a_t \sim \pi_{old}(\cdot|s_t)} \left[\frac{\pi(a_t|s_t)}{\pi_{old}(a_t|s_t)} \gamma^t A^{\pi_{old}}(s_t, a_t) \right] \right] \quad (\text{importance sampling}) \\ &\geq L_{\pi_{old}}(\pi) - C \max_s D_{KL}(\pi_{old}(\cdot|s) \parallel \pi(\cdot|s)) \quad (\text{surrogate objective})\end{aligned}$$

¹where $C = 4\epsilon\gamma/(1 - \gamma)$, $\epsilon = \max_{s,a} |A_{\pi}(s, a)|$

¹Schulman, John, et al. "Trust region policy optimization." *International conference on machine learning*. PMLR, 2015.

TRPO-Review : 2) Approximate trust region


$$\max_{\theta} L_{\theta_{old}}(\theta) - C \max_s D_{KL}(\theta_{old} \parallel \theta)$$

too small stepsize

TRPO-Review : 2) Approximate trust region

$$\max_{\theta} L_{\theta_{old}}(\theta) - C \max_s D_{KL}(\theta_{old} \parallel \theta)$$

too small stepsize




$$\begin{aligned} &\underset{\theta}{\text{maximize}} L_{\theta_{old}}(\theta) \\ &\text{subject to } D_{KL}^{\max}(\theta_{old}, \theta) \leq \delta. \end{aligned}$$

trust region constraint

TRPO-Review : 2) Approximate trust region

$$\max_{\theta} L_{\theta_{old}}(\theta) - C \max_s D_{KL}(\theta_{old} \parallel \theta)$$

too small stepsize



$$\begin{aligned} &\underset{\theta}{\text{maximize}} L_{\theta_{old}}(\theta) \\ &\text{subject to } D_{KL}^{\max}(\theta_{old}, \theta) \leq \delta. \end{aligned}$$


trust region constraint

bounded at every point in the state space

TRPO-Review : 2) Approximate trust region

$$\max_{\theta} L_{\theta_{old}}(\theta) - C \max_s D_{KL}(\theta_{old} \parallel \theta)$$

too small stepsize



$$\begin{aligned} &\underset{\theta}{\text{maximize}} L_{\theta_{old}}(\theta) \\ &\text{subject to } D_{KL}^{\max}(\theta_{old}, \theta) \leq \delta. \end{aligned}$$

trust region constraint



$$\begin{aligned} &\max_{\theta} L_{\theta_{old}}(\theta) \\ &\text{s.t. } \underbrace{\mathbb{E}^{\pi_{\theta_{old}}} [D_{KL}(\pi_{\theta_{old}} \parallel \pi_{\theta})]}_{\text{approximate trust region}} \leq \delta \end{aligned}$$

approximation (max. $\rightarrow$ avg.)

TRPO-Implementation

TRPO-Implementation

What info do we need when we implement **on-policy algorithms**?

1. (s_j, a_j, s'_j, r_j)
2. generalized advantage estimation(GAE)
3. MC estimates of value $V^{\pi_\phi}(s)$
4. probability $\pi_{\phi_{\text{old}}}(a_j|s_j)$
5. policy & value update

TRPO-Implementation

Generalized Advantage Estimation

policy gradient (with importance sampling)

$$\nabla_{\phi} J(\phi) = \mathbb{E}_{s \sim \rho_{\phi_{\text{old}}}, a \sim \pi_{\phi_{\text{old}}}} \left(A^{\pi_{\phi}}(s, a) \frac{\nabla_{\phi} \pi_{\phi}(a|s)}{\pi_{\phi_{\text{old}}}(a|s)} \right)$$

To estimate $A^{\pi_{\phi}}$, we use a separate value function approximator $V(s; \theta)$, and apply **generalized advantage estimation(GAE)**!

Furthermore, we use a stochastic policy $\pi_{\phi}(a|s)$ (Gaussian in our case).

GAE?

Given a trajectory $(s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T)$ generated by executing the current policy, we first compute

$$\delta_t = r_t + \gamma V(s_{t+1}; \theta) - V(s_t; \theta)$$

Then, GAE is computed as follows:


$$\hat{A}(s_t, a_t) = \sum_{l=0}^{\infty} (\lambda \gamma)^l \delta_{t+l}$$

TRPO-Implementation

Generalized Advantage Estimation


- $\lambda = 0$

$$\hat{A}(s_t, a_t) = r_t + \gamma V(s_{t+1}; \theta) - V(s_t; \theta)$$



- $\lambda = 1$

$$\hat{A}(s_t, a_t) = r_t + r_{t+1} + r_{t+2} + \dots$$



training $V(s; \theta)$?

given a trajectory $(s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T)$ generated from π_ϕ ,
we can compute Monte-Carlo estimates of $V(s_0), V(s_1), \dots, V(s_T)$
as follows:

$$V(s_t) = \sum_{\tau=t}^{T-1} \gamma^{\tau-t} r_\tau + \gamma^T V(s_T)$$

This is the **target** for value function update!

TRPO-Implementation

Value estimate

```
self._obs_mem = np.zeros(shape=(lim, dimS))  
self._act_mem = np.zeros(shape=(lim, dimA))  
self._rew_mem = np.zeros(shape=(lim,))  
self._val_mem = np.zeros(shape=(lim,))  
self._log_prob_mem = np.zeros(shape=(lim,))
```

collected during agent-env interaction

```
# memory of cumulative rewards which are MC-estimates of the current value function  
self._target_v_mem = np.zeros(shape=(lim,))  
# memory of GAE( $\lambda$ )-estimate of the current advantage function  
self._adv_mem = np.zeros(shape=(lim,))
```

computed when sampling a single episode is done

TRPO-Implementation

Approximate KL divergence

$$\begin{array}{ll} \max_{\theta} & L_{\theta_{old}}(\theta) \\ \text{s.t.} & \underbrace{\mathbb{E}^{\pi_{\theta_{old}}}[D_{KL}(\pi_{\theta_{old}} \parallel \pi_{\theta})]}_{\text{approximate trust region}} \leq \delta \end{array}$$

Taylor approximation

$$\begin{array}{ll} \max_{\theta} & g^{\top}(\theta - \theta_{old}) \\ \text{s.t.} & \frac{1}{2}(\theta - \theta_{old})^{\top} H(\theta - \theta_{old}) \leq \delta \end{array}$$

obj. $\rightarrow$ first-order

const. $\rightarrow$ second-order

where $H := \nabla_{\theta}^2[D_{KL}(\pi_{\theta_{old}} \parallel \pi_{\theta})]$,

$g := \nabla_{\theta} L_{\theta_{old}}(\theta)|_{\theta=\theta_{old}}$

TRPO-Implementation

Approximate KL divergence

$$\begin{aligned} \max_{\theta} \quad & L_{\theta_{old}}(\theta) \\ \text{s.t.} \quad & \underbrace{\mathbb{E}^{\pi_{\theta_{old}}} [D_{KL}(\pi_{\theta_{old}} \parallel \pi_{\theta})]}_{\text{approximate trust region}} \leq \delta \end{aligned}$$

Taylor approximation

$$\begin{aligned} \max_{\theta} \quad & g^{\top}(\theta - \theta_{old}) \\ \text{s.t.} \quad & \frac{1}{2}(\theta - \theta_{old})^{\top} H(\theta - \theta_{old}) \leq \delta \end{aligned}$$

Solution.

$$\theta_{new} = \theta_{old} + \sqrt{\frac{2\delta}{g^{\top} H^{-1} g}} H^{-1} g$$

TRPO-Policy & Value Update

Unfortunately, the parameter updates in TRPO...

1. Fisher-vector product
2. conjugate gradient descent $\rightarrow H^{-1}g$
3. backtracking line search $\rightarrow$ still trust region?

Solution.

$$\theta_{new} = \theta_{old} + \sqrt{\frac{2\delta}{g^\top H^{-1}g}} H^{-1}g$$

TRPO-Implementation

Policy & Value Update

1. Fisher-vector product

```
1 def fisher_vector_product(v, actor, obs_batch, cg_damping=1e-2):
2     v.detach()
3     kl = torch.mean(kl_div(actor=actor, old_actor=actor, obs_batch=obs_batch))
4     kl_grads = torch.autograd.grad(kl, actor.parameters(), create_graph=True)
5     kl_grad = torch.cat([grad.view(-1) for grad in kl_grads])
6     kl_grad_p = torch.sum(kl_grad * v)
7     Iv = torch.autograd.grad(kl_grad_p, actor.parameters()) # product of Fisher information I and v
8     Iv = flatten(Iv)
9     return Iv + v * cg_damping
```

Basically, Fish information matrix is **Hessian** of a function.

⇒ costly (requires n^2 numbers of 2nd-order differentiation)

However, computation of Ax can easily be done using automatic differentiation:
use

$$(\nabla_{\theta}^2 f)v = \nabla_{\theta}(\underbrace{\nabla_{\theta} f \cdot v}_{\text{scalar}}).$$

TRPO-Implementation

Policy & Value Update

2. conjugate gradient method

```
1  def cg(f_Ax, b, actor, obs_batch, cg_iters=10, residual_tol=1e-10):
2      p = b.clone()
3      r = b.clone()
4      x = torch.zeros_like(b)
5      rdotr = r.dot(r)
6      for i in range(cg_iters):
7          z = f_Ax(p, actor, obs_batch)
8          v = rdotr / p.dot(z)
9          x += v*p
10         r -= v*z
11         newrdotr = r.dot(r)
12         mu = newrdotr/rdotr
13         p = r + mu*p
14         rdotr = newrdotr
15         if rdotr < residual_tol:
16             break
17     return x
```

What is conjugate gradient for?

⇒ To efficiently solve $As = g$

- Solve $\frac{1}{2}s^T H s = g^T s$, rather than $Hs = g$
- Search along the conjugate directions

TRPO-Implementation

Policy & Value Update

3. backtracking line search

```
1 def backtracking_line_search(...):
2     backtrac_coef = 1.0
3     alpha = 0.5
4     beta = 0.5
5     flag = False
6     expected_improve = (actor_loss_grad * maximal_step).sum(0, keepdim=True)
7     for i in range(10):
8         new_params = params + backtrac_coef * maximal_step
9         update_model(actor, new_params)
10        new_actor_loss = surrogate_loss(actor, adv, states, old_policy.detach(), actions)
11        loss_improve = new_actor_loss - actor_loss
12        expected_improve *= backtrac_coef
13        improve_condition = loss_improve / expected_improve
14        kl = kl_div(actor=actor, old_actor=old_actor, obs_batch=states)
15        kl = kl.mean()
16        if kl < max_kl and improve_condition > alpha:
17            flag = True
18            break
19        backtrac_coef *= beta
20    if not flag:
21        params = flat_params(old_actor)
22        update_model(actor, params)
```

$\theta \leftarrow \theta + \beta s$ where s : search direction found through CG

check if new policy stays within trust region

decrease β and repeat the verification

TRPO-Implementation

Policy & Value Update

3. backtracking line search

```
1 def backtracking_line_search(...):
2     backtrac_coef = 1.0
3     alpha = 0.5
4     beta = 0.5
5     flag = False
6     expected_improve = (actor_loss_grad * maximal_step).sum(0, keepdim=True)
7     for i in range(10):
8         new_params = params + backtrac_coef * maximal_step
9         update_model(actor, new_params)
10        new_actor_loss = surrogate_loss(actor, adv, states, old_policy.detach(), actions)
11        loss_improve = new_actor_loss - actor_loss
12        expected_improve *= backtrac_coef
13        improve_condition = loss_improve / expected_improve
14        kl = kl_div(actor=actor, old_actor=old_actor, obs_batch=states)
15        kl = kl.mean()
16        if kl < max_kl and improve_condition > alpha:
17            flag = True
18            break
19        backtrac_coef *= beta
20    if not flag:
21        params = flat_params(old_actor)
22        update_model(actor, params)
```

$\theta \leftarrow \theta + \beta s$ where s : search direction found through CG

check if new policy stays within trust region

decrease β and repeat the verification

TRPO-Implementation

Policy & Value Update

3. backtracking line search

```
1 def backtracking_line_search(...):
2     backtrac_coef = 1.0
3     alpha = 0.5
4     beta = 0.5
5     flag = False
6     expected_improve = (actor_loss_grad * maximal_step).sum(0, keepdim=True)
7     for i in range(10):
8         new_params = params + backtrac_coef * maximal_step
9         update_model(actor, new_params)
10        new_actor_loss = surrogate_loss(actor, adv, states, old_policy.detach(), actions)
11        loss_improve = new_actor_loss - actor_loss
12        expected_improve *= backtrac_coef
13        improve_condition = loss_improve / expected_improve
14        kl = kl_div(actor=actor, old_actor=old_actor, obs_batch=states)
15        kl = kl.mean()
16        if kl < max_kl and improve_condition > alpha:
17            flag = True
18            break
19        backtrac_coef *= beta
20    if not flag:
21        params = flat_params(old_actor)
22        update_model(actor, params)
```

$\theta \leftarrow \theta + \beta s$ where s : search direction found through CG

check if new policy stays within trust region

decrease β and repeat the verification

TRPO-Implementation

Policy & Value Update

getting together..

```
g = torch.cat([grad.view(-1) for grad in loss_grad]).data
s = cg(fisher_vector_product, g, agent.pi, states)

sAs = torch.sum(fisher_vector_product(s, agent.pi, states) * s, dim=0, keepdim=True)
step_size = torch.sqrt(2 * delta / sAs)[0]
step = step_size * s
```

$$\theta_{new} = \theta_{old} + \sqrt{\frac{2\delta}{g^\top H^{-1}g}} H^{-1}g$$

TRPO-Implementation

Policy & Value Update

getting together..

```
g = torch.cat([grad.view(-1) for grad in loss_grad]).data
s = cg(fisher_vector_product, g, agent.pi, states)

sAs = torch.sum(fisher_vector_product(s, agent.pi, states) * s, dim=0, keepdim=True)
step_size = torch.sqrt(2 * delta / sAs)[0]
step = step_size * s
```

Why?

- $s \approx H^{-1}g$
- The denominator aims to measure $s^\top H s$ with s being used in the algorithm rather than the theoretical $H^{-1}g$.


TRPO- Complete Implementation

TRPO-Network Architecture


Key components of TRPO agent:

- policy network
 - can handle both discrete/continuous actions
 - network output : **probability distribution** over action space (cf. DDPG)

ex) discrete action space



ex) continuous action space



TRPO-Network Architecture

```
1 class Actor(nn.Module):
2     def __init__(self, obs_dim, act_dim, hidden1, hidden2):
3         super(Actor, self).__init__()
4         self.fc1 = nn.Linear(obs_dim, hidden1)
5         self.fc2 = nn.Linear(hidden1, hidden2)
6         self.fc3 = nn.Linear(hidden2, act_dim) # for \mu
7         self.fc4 = nn.Linear(hidden2, act_dim) # for \sigma
8
9     def forward(self, obs):
10        x = torch.tanh(self.fc1(obs))
11        x = torch.tanh(self.fc2(x))
12        mu = self.fc3(x)
13        log_sigma = self.fc4(x)
14        sigma = torch.exp(log_sigma)
15        return mu, sigma
16
17    def log_prob(self, obs, act):
18        mu, sigma = self.forward(obs)
19        act_distribution = Independent(Normal(mu, sigma), 1)
20        log_prob = act_distribution.log_prob(act)
21        return log_prob
```

1. Actor Network (By actor we just mean policy)

Why $\log \sigma(s)$ instead of $\sigma(s)$?
 $\implies$ We have a constraint $\sigma(s) > 0$!

gives $\log \pi(a^i | s^i)$'s for (s^i, a^i) pairs

TRPO-Network Architecture

```
1 class Critic(nn.Module):
2     # critic V(s ; \theta)
3     def __init__(self, obs_dim, hidden1, hidden2):
4         super(Critic, self).__init__()
5         self.fc1 = nn.Linear(obs_dim, hidden1)
6         self.fc2 = nn.Linear(hidden1, hidden2)
7         self.fc3 = nn.Linear(hidden2, 1)
8
9     def forward(self, obs):
10         x = torch.tanh(self.fc1(obs))
11         x = torch.tanh(self.fc2(x))
12
13         return self.fc3(x)
```

2. Critic Network : super-easy!

TRPO-Action Selection by Agent

```
1 class TRPOAgent:
2     def __init__(self, obs_dim, act_dim, hidden1=64, hidden2=32):
3         ...
4
5     def act(self, obs):
6         obs = torch.tensor(obs, dtype=torch.float).to(device)
7         with torch.no_grad():
8             mu, sigma = self.pi(obs)
9             act_distribution = Independent(Normal(mu, sigma), 1)
10            action = act_distribution.sample()
11            log_prob = act_distribution.log_prob(action)
12            val = self.V(obs)
13        action = action.cpu().numpy()
14        log_prob = log_prob.cpu().numpy()
15        val = val.cpu().numpy()
16
17        return action, log_prob, val
```

Gaussian distribution $\mathcal{N}(\mu(s^i), \sigma(s^i))$

$a^i \sim \pi_\phi(s^i) := \mathcal{N}(\mu(s^i), \sigma(s^i))$

$\log \pi_\phi(a^i | s^i)$

$V(s^i; \theta)$

TRPO-Complete Implementation

```
1 def train(env, agent, max_iter, gamma=0.99, lr=3e-4, lam=0.95, delta=1e-3, steps_per_epoch=4000, eval_interval=4000):
2     ...
3     for epoch in range(num_epochs):
4         state = env.reset()
5         step_count = 0
6         ep_reward = 0
7         for t in range(steps_per_epoch):
8             action, log_prob, v = agent.act(state)
9             next_state, reward, done, _ = env.step(action)
10            memory.append(state, action, reward, v, log_prob)
11            ep_reward += reward
12            step_count += 1
13            if (step_count == max_ep_len) or (t == steps_per_epoch - 1):
14                s_last = torch.tensor(next_state, dtype=torch.float).to(device)
15                v_last = agent.V(s_last).item()
16                memory.compute_values(v_last)
17            elif done:
18                v_last = 0.0
19                memory.compute_values(v_last)
```

TRPO-Complete Implementation

```
21         state = next_state
22         if done:
23             state = env.reset()
24             step_count = 0
25             ep_reward = 0
26             total_t += 1
27         update(agent, memory, critic_optim, delta, num_updates=1)
28     return
```

PPO - Review

Recap : TRPO

$$\begin{array}{ll} \max_{\theta} & L_{\theta_{old}}(\theta) \\ \text{s.t.} & \underbrace{\mathbb{E}^{\pi_{\theta_{old}}} [D_{KL}(\pi_{\theta_{old}} \parallel \pi_{\theta})]}_{\text{approximate trust region}} \leq \delta \end{array}$$

$$L_{\theta_{old}}(\theta) := \mathbb{E}_t \left[\underbrace{\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}}_{:= r_t(\theta)} \gamma^t A^{\pi_{\theta_{old}}}(s_t, a_t) \right]$$

Proximal Policy Optimization : Idea

$$\begin{array}{ll} \max_{\theta} & L_{\theta_{old}}(\theta) \\ \text{s.t.} & \underbrace{\mathbb{E}^{\pi_{\theta_{old}}} [D_{KL}(\pi_{\theta_{old}} \parallel \pi_{\theta})]}_{\text{approximate trust region}} \leq \delta \end{array}$$

$$L_{\theta_{old}}(\theta) := \mathbb{E}_t \left[\underbrace{\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}}_{:= r_t(\theta)} \gamma^t A^{\pi_{\theta_{old}}}(s_t, a_t) \right]$$

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$

Proximal Policy Optimization : Idea

$$\begin{array}{ll} \max_{\theta} & L_{\theta_{old}}(\theta) \\ \text{s.t.} & \underbrace{\mathbb{E}^{\pi_{\theta_{old}}} [D_{KL}(\pi_{\theta_{old}} \parallel \pi_{\theta})]}_{\text{approximate trust region}} \leq \delta \end{array}$$

$$L_{\theta_{old}}(\theta) := \mathbb{E}_t \left[\underbrace{\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}}_{:= r_t(\theta)} \gamma^t A^{\pi_{\theta_{old}}}(s_t, a_t) \right]$$

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$

removes the incentive for moving ratio outside of the interval $[1 - \epsilon, 1 + \epsilon]$

$\approx$ KL penalty

Proximal Policy Optimization(PPO) : What's different?

```
1  for i in range(num_iter):
2      log_probs, ent = self.pi.compute_log_prob(states, actions)
3      r = torch.exp(log_probs - old_log_probs)
4      clipped_r = torch.clamp(r, 1 - self.epsilon, 1 + self.epsilon)
5      single_step_obj = torch.min(r * A, clipped_r * A)
6      pi_loss = -torch.mean(single_step_obj)
7      v = self.V(states)
8      V_loss = torch.mean((v - target_v) ** 2)
9      ent_bonus = torch.mean(ent)
10     loss = pi_loss + 0.5 * V_loss - 0.01 * ent_bonus
11     self.optimizer.zero_grad()
12     loss.backward()
13     self.optimizer.step()
14     if i == num_iter - 1:
15         kl = torch.mean(old_log_probs - log_probs).item()
16         if kl > 1.5 * self.kl_threshold:
17             break
18     return
```

simple, surrogate loss function for policy network parameters!
⇒ 1st-order optimization (no Hessian needed)

Thank you!